Alexander Panfilov

kotekjedi@gmail.com | | | | Tübingen, Germany

EDUCATION

IMPRS-IS / ELLIS | PhD Machine Learning
May 2024 – Aug 2027 | Tübingen, Germany

University of Tübingen | MSc Machine Learning
Oct 2021 – Apr 2024 | Tübingen, Germany

ITMO University | BSc Software Engineering
Sep 2017 – Jun 2021 | Saint Petersburg, Russia

RESEARCH EXPERIENCE

Google DeepMind | Research Scientist
Nov 2026 – TBD | London, United Kingdom

Frontier Risks Mitigations Team (Gemini Safety)

MATS 9.0/9.1 | Research Fellow
Jan 2026 – Sep 2026 | London, United Kingdom

GDM Stream (Zimmermann, Jenner, Lindner)

- Benchmarked models' capability to detect safety interventions in AI control pipelines, with results accepted at the **ICLR 2026 Workshop (oral)** on Catch, Adapt, and Operate: Monitoring ML Models Under Drift.

ELLIS Institute Tübingen / MPI for Intelligent Systems | Doctoral Student
May 2024 – Nov 2026 | Tübingen, Germany

Jonas Geiping's Group (SEAL), co-supervised by Maksym Andriushchenko

- Worked on topics spanning ML security and AI safety, including jailbreaking, prompt-injection defenses, AI control, and deception in LLMs, with **6 papers accepted at top ML conferences**: 1 at **ICML**, 4 at **ICLR**, and 1 at **NeurIPS**.

University of Tübingen / MPI for Intelligent Systems | Research Assistant
May 2022 – Apr 2024 | Tübingen, Germany

Wieland Brendel's Group (RobustML)

- Proposed and implemented a novel regularization method enabling combinatorial generalization in object-centric models. Results accepted at **ICLR 2024 (oral)**, ranking in the **top 1.2%** of submitted papers.

ITMO University | Research Assistant
Nov 2019 – Oct 2021 | Saint Petersburg, Russia

Machine Learning Lab

- Contributed to research on exploiting simplicity bias for adversarial training and domain adaptation via optimal transport, with results published at the **ICML 2022 Workshop** on AdvML Frontiers and **CoLLAs 2023**.

Center for Learning Analytics

- Developed digital student profiles and engineered a machine learning system to predict academic outcomes and potential student expulsions aiding early intervention planning.

Robert Bosch | Research Intern
Jun 2020 – Nov 2020 | Saint Petersburg, Russia

- Worked on disentanglement in VAEs for identifying cost-effective yet high-performing motor designs.

SELECTED PUBLICATIONS

- preprint – equal contribution
  Stealing Reasoning Traces from Proprietary LLM APIs
  Panfilov, A.* , Schmotz, D.*, Shumailov, I.*, Schaeffer, J.*, Beurer-Kellner, L., Prabhu, A., Geiping, J., and Andriushchenko, M.
  arXiv preprint arXiv:2608.09867, 2026

- conference paper – equal contribution
  Claudini: Autoresearch Discovers State-of-the-Art Adversarial Attack Algorithms for LLMs
  Panfilov, A.* , Romov, P.*, Shilov, I.*, de Montjoye, Y.-A., Geiping, J., and Andriushchenko, M.
  NeurIPS 2026

- conference paper – equal contribution
  Adaptive Attacks on Trusted Monitors Subvert AI Control Protocols
  Terekhov, M.*, Panfilov, A.* , Dzenhaliou, D.*, Gulcehre, C., Andriushchenko, M., Prabhu, A., and Geiping, J.
  ICLR 2026

- conference paper – equal contribution
  Strategic Dishonesty Can Undermine AI Safety Evaluations of Frontier LLMs
  Panfilov, A.* , Kortukov, E.*, Nikolić, K., Bethge, M., Lapuschkin, S., Samek, W., Prabhu, A., Andriushchenko, M., and Geiping, J.
  ICLR 2026

- conference paper
  Capability-Based Scaling Trends for LLM-Based Red-Teaming
  Panfilov, A. , Kassianik, P., Andriushchenko, M., and Geiping, J.
  ICLR 2026

- conference paper – equal contribution
  An Interpretable N-gram Perplexity Threat Model for Large Language Model Jailbreaks
  Boreiko, V.*, Panfilov, A.* , Voracek, V., Hein, M., and Geiping, J.
  ICML 2025

ACADEMIC SERVICE

- **Co-advising with Maksym Andriushchenko**:
  - Lena Libon (ETH Zürich, Master's thesis) — Probe-Based Post-Training of LLMs (ICML 2026 Workshop paper)
  - Derck Prinzhorn (University of Amsterdam, Master's thesis) — Model Organisms-Based Red-Teaming for AI Control
- **Program Committee**: AISec 2026, NeurIPS 2026 Workshop on Foundations of Language Model Security (FLMSec)
- **Reviewer**: NeurIPS 2025, NeurIPS 2026, ICLR 2026, ICLR 2027, AISec 2026, ICLR 2026 Trustworthy AI Workshop, ICLR 2026 CAO Workshop, ICML 2026 AI4Research Workshop, NeurIPS 2026 FLMSec Workshop
- **Volunteer**: LLM Safety and Security Workshop, ELLIS UnConference 2025 • Foundations of Language Model Security Workshop, EurIPS 2025
- Examiner for two MSc theses (ITMO University 2023, HSE St. Petersburg 2023)

INVITED TALKS

- **Commission nationale de l'informatique et des libertés (CNIL)**: France's independent public authority for data protection (remote) — Oct 09, 2026
- **Cohere Labs Community**: invited talk, Privacy, Security & Policy group (remote) — Oct 01, 2026
- **MIT CSAIL**: invited talk, NLP Seminar (remote) — Sep 09, 2026
- **ETH Zürich**: invited talk with Florian Tramèr's SPY Lab, SRI Lab, and Meta researchers — Sep 03, 2026
- **MATS 10.0 Symposium**: keynote (London) — Aug 21, 2026
- **UK AI Security Institute**: UK government AI safety research institute (invited talk, London) — Aug 20, 2026
- **OpenAI**: research presentation (remote) — Aug 15, 2026
- **Google DeepMind**: research presentation, Frontier Risks Mitigations team (Gemini Safety; remote) — Jun 19, 2026
- **Google DeepMind and MATS**: invited talk, AGI Safety Team (London) — Jun 18, 2026
- **Imperial College London**: invited talk, Yves-Alexandre de Montjoye's group seminar (London) — Feb 09, 2026
- **Google**: invited talk, ML Red Teaming Seminar (remote) — Jun 23, 2025
- **EPFL**: invited talk, Nicolas Flammarion's group seminar (Lausanne) — Nov 05, 2024

INTERVIEWS & PRESS COMMENTS

- **WIRED**: A New Trick Reveals AI Models' Inner Thoughts.
- **Scientific American**: AI companies say rivals are 'distilling' their models. What does that even mean?
- **The Stack**: Stolen valor? How researchers discovered some open-weight models might have cut corners.
- **Machine Learning Street Talk**: Why Frontier AI Labs Fight to Hide Chain of Thought — Ilia Shumailov & Alexander Panfilov.

AWARDS & SCHOLARSHIPS

- **Apart Research ARENA 6.0 Mechanistic Interpretability Hackathon (2025), 3rd place**: Awarded for the project investigating whether white-box deception detectors catch deception or instruction to deceive.
- **ELSA Grant (2025)**: Awarded a research travel grant from European Lighthouse on Secure and Safe AI (∼ 3k euro).
- **DAAD Scholarship (2021), Top 3%**: Selected as one of ∼30 Russian students from ∼1,000 applicants for a two-year DAAD-funded master's program in Germany (∼ 30k euro).
- **"Ya-Professional" Student Olympiad Winner (2021), Top 2%**: Achieved prizeholder status in AI and ML tracks, with only 3,881 out of 177,100 participants (among all tracks) receiving this distinction.

INDUSTRY EXPERIENCE

X5 Group | Data Scientist

Nov 2020 – Oct 2021 | Moscow, Russia (remote)

- Designed and conducted A/B testing experiments to evaluate the efficacy of various business initiatives, performed ad-hoc analytics to support decision-making processes within Russia's largest offline retail chain.
- Mentored three interns, all subsequently securing full-time roles within the company.

Yandex | Machine Learning Engineer Intern

Feb 2020 – May 2020 | Moscow, Russia

- Optimized the push notification system at Yandex.Zen for personalized timing of notifications.

FULL PUBLICATIONS LIST

- preprint
  Alignment via Training Against Probes Without Losing Monitorability
  Libon, L., Panfilov, A. , Rank, B., Chen, X., Geiping, J., and Andriushchenko, M.
  arXiv preprint arXiv:2609.38645, 2026

- preprint
  Distillation Defenses Easily Break After Reinforcement Learning
  Javaheri, S., Panfilov, A. , Britton, O., Gal, Y., and Gideoni, Y.
  arXiv preprint arXiv:2609.35699, 2026

- preprint – equal contribution
  Stealing Reasoning Traces from Proprietary LLM APIs
  Panfilov, A.* , Schmotz, D.*, Shumailov, I.*, Schaeffer, J.*, Beurer-Kellner, L., Prabhu, A., Geiping, J., and Andriushchenko, M.
  arXiv preprint arXiv:2608.09867, 2026

- preprint – equal contribution
  CIAware-Bench: Benchmarking Control Intervention Awareness Across Frontier LLMs
  Schaeffer, J.*, Panfilov, A.* , Jiralerspong, T.*, Lajoie, G., Geiping, J., Bengio, Y., and Zimmermann, R. S.
  arXiv preprint arXiv:2606.11063, 2026

- workshop paper
  Fine-tuning with Harmfulness Probes Leads to Natural Refusals
  Libon, L., Panfilov, A. , Rank, B., Geiping, J., and Andriushchenko, M.
  ICML 2026 Mechanistic Interpretability Workshop

- conference paper – equal contribution
  Claudini: Autoresearch Discovers State-of-the-Art Adversarial Attack Algorithms for LLMs
  Panfilov, A.* , Romov, P.*, Shilov, I.*, de Montjoye, Y.-A., Geiping, J., and Andriushchenko, M.
  NeurIPS 2026

- workshop paper – equal contribution
  Measuring Control Intervention Awareness Across Frontier LLMs
  Schaeffer, J.*, Jiralerspong, T.*, Panfilov, A.* , and Zimmermann, R. S.
  ICLR 2026 Catch, Adapt, and Operate Workshop (Oral)

- conference paper – equal contribution
  Adaptive Attacks on Trusted Monitors Subvert AI Control Protocols
  Terekhov, M.*, Panfilov, A.* , Dzenhaliou, D.*, Gulcehre, C., Andriushchenko, M., Prabhu, A., and Geiping, J.
  ICLR 2026

- conference paper – equal contribution
  Strategic Dishonesty Can Undermine AI Safety Evaluations of Frontier LLMs
  Panfilov, A.* , Kortukov, E.*, Nikolić, K., Bethge, M., Lapuschkin, S., Samek, W., Prabhu, A., Andriushchenko, M., and Geiping, J.
  ICLR 2026

- conference paper
  Capability-Based Scaling Trends for LLM-Based Red-Teaming
  Panfilov, A. , Kassianik, P., Andriushchenko, M., and Geiping, J.
  ICLR 2026; also in *ICML 2025 Workshop on Reliable and Responsible Foundation Models*

- conference paper
  ASIDE: Architectural Separation of Instructions and Data in Language Models
  Zverev E., Kortukov E., Panfilov, A. , Volkova A., Tabesh S., Lapuschkin S., Samek W., and Lampert H Ch.
  ICLR 2026; also in *ICLR 2025 Workshop on Building Trust in Language Models and Applications (Oral)*

- conference paper – equal contribution
  An Interpretable N-gram Perplexity Threat Model for Large Language Model Jailbreaks
  Boreiko, V.*, Panfilov, A.* , Voracek, V., Hein, M., and Geiping, J.
  ICML 2025; also in *NeurIPS 2024 Red Teaming Gen AI Workshop (Oral)*

- conference paper – equal contribution
  Provable Compositional Generalization for Object-Centric Learning
  Wiedemer, T.*, Brady, J.*, Panfilov, A.* , Juhos, A.*, Bethge, M., and Brendel, W.
  ICLR 2024 (Oral)

- conference paper
  A Minimalist Approach for Domain Adaptation with Optimal Transport
  Asadulaev, A., Shutov, V., Korotin, A., Panfilov, A. , Kontsevaya, V., and Filchenkov, A.
  Proceedings of The 2nd Conference on Lifelong Learning Agents, PMLR 232:1009-1024, 2023

- workshop paper
  Easy Batch Normalization
  Asadulaev, A., Panfilov, A. , and Filchenkov, A.
  ICML 2022 AdvML Frontiers Workshop, 2022

- workshop paper
  Multi-step domain adaptation by adversarial attack to $\mathcal{H}\Delta\mathcal{H}$ -divergence
  Asadulaev, A., Panfilov, A. , and Filchenkov, A.
  ICML 2022 AdvML Frontiers Workshop, 2022

- conference paper
  Recommender system for an academic supervisor with a matrix normalization approach
  Kazakovtsev, V., Oreshin, S., Serdyukov, A., Krasheninnikov, E., Muravyov, S., Bezvinnyi, A., Panfilov, A. , Glukhov, I., Kaliberda, Y., Masalskiy, D., Podolenchuk, T., and Khlopotov, M.
  Proceedings of The 2020 1st International Conference on Control, Robotics and Intelligent System (CCRIS '20)

- conference paper
  Implementing a Machine Learning Approach to Predicting Students' Academic Outcomes
  Oreshin, S., Filchenkov, A., Petrusha, P., Krasheninnikov, E., Panfilov, A. , Glukhov, I., Kaliberda, Y., Masalskiy, D., Serdyukov, A., Kazakovtsev, V., Khlopotov, M., Podolenchuk, T., Smetannikov, I., and Kozlova, D.
  Proceedings of The 2020 1st International Conference on Control, Robotics and Intelligent System (CCRIS '20)

- chapter
  The Use of Students' Digital Portraits in Creating Smart Higher Education: A Case Study of the AI Benefits in Analyzing Educational and Social Media Data
  Oreshin, S., Filchenkov, A., Kozlova, D., Petrusha, P., Panfilov, A. , Glukhov, I., Krasheninnikov, E. and Buraya, I.
  In: Uskov, V., Howlett, R., Jain, L. (eds) Smart Education and e-Learning 2020